Human Activity Recognition Using Accelerometer Data

Aishwarya G Shenoy[1], **Rohit Keswani**[2], **Subhadeep Das**[3]

Abstract - **Human Activity Recognition (HAR) has a wide range of applications due to the widespread usage of acquisition devices such as smartphones and its ability to capture human activity data. The ability to retrieve deeply embedded information for precise detection and its interpretation has been transformed by breakthroughs in Artificial Intelligence (AI). In this paper, the time series dataset, acquired from Wireless Sensor Data Mining Lab (WISDM) Lab, is used to extract features of common human activities from a raw signal data of smartphone accelerometer. A 2D convolutional neural network is used to visualize the data.**

Keywords: Human Activity Recognition, StratifiedK-fold cross validation, CNN

I. INTRODUCTION

Human Activity Recognition (HAR) is the technique of utilizing Artificial Intelligence (AI) to recognize and classify human activities from raw activity data collected by a range of devices. A few examples of such devices include smartphones and smartwatches. Smartphones and smartwatches consist of accelerometer sensors (tri-axial accelerometers) that is used to measure acceleration in three different dimensions. The orientation of the device can be determined by these accelerometers, which can be helpful information for activity detection. The time series dataset used in this paper consists of raw data that is collected from 36 different participants performing different activities such as walking, jogging, sitting, standing, ascending and descending for specific time periods. The data is collected using a sample rate of 20 Hz (1 sample every 50 milliseconds) which is equivalent to 20 samples per second. There are six different attributes in the dataset namely user, activity, timestamp, x-axis, y-axis and z-axis. The dataset contains over a million rows and 6 columns that needs to be cleaned and processed. The challenge of significant feature extraction from the raw sensor data has become easier with the introduction of Deep Learning (DL) in the HAR domain. A 2D Convolutional Neural Network (CNN) is used to classify the data and a StratifiedKfold cross validation is used to split the data into train and test data.

II. APPROACH

A. Pre-processing and Data Exploration

The raw dataset, containing over a million rows, has a lot of noise that needs to be eliminated. In order to improve the accuracy, we first clean the dataset by condensing the dataframe. We are shortening the dataframe by reducing the number of participants to five. Then, the existence of any null values is checked. After checking for null values, we determine the data distribution of five participants performing several activities. Through this data distribution we can deduce that the data is unbalanced. Standing examples are fewer in number as compared to Walking examples. As a result, using this data directly will cause overfitting, which will skew the results in favour of walking. To avoid this issue, we need to balance the data by eliminating unnecessary data that is covered in the further section of this paper. Now, we plot the tri-axial orientation to observe the variation in the data and convert all the data from string type to float type. Figure 1 represents the time taken to perform different activities by the participants.


Fig. 1. Representation of data in a bar graph

In order to determine the ratio between training and test set, the variation in the signal values of time taken by the participants while engaging in different activities has to be observed. Firstly, let us visualize the time series signal waves of a single participant. In this scenario, we have selected participant 5 as an example. The following figures represent the time taken by participant 5 to perform activities such as Walking, Jogging, Descending Downstairs, Ascending Upstairs, Sitting and Standing in tri-axial orientation.

From the figures, we can deduce that the signals show a drastic change in the period behaviour for the activities such as Walking, Jogging, Upstairs and Downstairs but there is a minimal movement while performing stationary activities such as Sitting and Standing.

Fig. 2. Walking

Fig. 3. Jogging


Fig. 4. Downstairs

Fig. 5. Upstairs

Fig. 6. Sitting

Fig. 7. Standing

B. Feature Extraction - Statistical Measures

The statistical measures such as standard deviation that measure the features of the signal windows are extracted to identify patterns in the data. We evaluate the standard deviation with respect to each user and the activity performed in three different dimensions (x axis, y axis and z axis). The mean of standard deviation across all users is evaluated and we have used a heatmap to visualize the results. The feature extraction has been performed on the dataset before standardization to visualize the unbalanced nature of the dataset. This helps us to eliminate unnecessary data values which would lead our model to be skewed after training.


Fig. 8. Heatmap

C. Standardization of Data

After visualizing the unbalanced dataset, the data is balanced by considering only the first 4339 lines for each activity and we create a balanced dataset. As the class Standing has the least samples (4339), we select the exact amount of samples for the other classes as well. We convert the tri-axial values in x, y and z columns from string type to numeric type. As standard classification algorithms cannot be applied directly, a sliding window technique is used to transform the raw series data. The data is first split into a window of 4 seconds and then we aggregate 80 raw samples present in each of the 4 second window to generate new features. The reason behind choosing a window size of 4 seconds is because a window size that is either too small or too large can affect how well the motion is captured in the modified dataset for training. We select the action that occurs most frequently in that window to apply a class label to the modified characteristics. The frequency is multiplied by 4 seconds resulting in a hop size of 40 but there will be some overlapping due to this. The overlapping ensures that every successive row in the transformed dataset contains some information from the data in previous window as well. Each of the six activities contains 4339 examples with a total of (4339*6) observations. This is equal to about 649 when divided by the hop size.

D. StratifiedKFold Cross Validation

Supervised learning algorithms aim to simulate how features (independent variables) and labels relate to each other. The model is first trained by using the features and labels of some observations. Then the model is tested by only providing the features and expecting it to predict labels from the given features. As a result, the data is split into training and testing subsets. The model learns from the training set and we can predict the accuracy using the testing set. In this scenario, a StratifiedKFold cross validation method is used to split the data for training and testing. The class distribution in the dataset is retained during the training and testing splits. After applying StratifiedKFold Cross Validation, the dataset is 3 dimensional. However each sample in the dataset is 2 dimensional and needs to reshaped as the Convolutional Neural Network (CNN) model uses 3 dimensional data.

III. RESULTS

A. 2D Convolutional Neural Network (CNN) Model

A 2D Convolutional Neural Network (CNN) comprises of convolutional layers that helps in improving the model performance and accuracy with the notable decrease in the computation time. The main challenge in building a model is to achieve a good balance between the computational time and accuracy that can be achieved using a CNN model. We have used a sequential model combined with a 2D convolution layer as it is appropriate for a simple stack of layers where each layer has one input and one output tensor. A tensor of outputs is produced by the 2D convolution layer by wrapping a convolution kernel with the input layers. In this scenario, a total of 16 filters with a size of (2,2) are assigned in the

first convolution layer. We have used ReLU, also known as the rectified linear activation function, that returns the input directly in the case of a positive outcome and returns zero if the result is negative. The outgoing edges of hidden units are randomly set to 0 at each update of the training phase using a dropout layer. The probability at which the outputs of the layers are dropped out is indicated by the value passed in dropout. The data is transformed into a 1 dimensional array using the flatten() function before proceeding to the next layer. The convolution model has a dense layer with 64 neurons and an output layer with 6 neurons for each of the 6 classes. The results are then produced as a probability distribution using softmax. Next, we compile and fit the model to the training data by using 10 epochs. At the end of each epoch, the validation of data is checked to evaluate the loss and model metrics.


Fig. 9. Model Loss curve

Fig. 10. Model Accuracy curve

The overall model accuracy and model loss is plotted. From the figures we can deduce that the accuracy of the training set

becomes better with the iteration of each epoch. Our model has learned from the training and is able to accurately predict each class in the validation set. Overall, we have achieved an accuracy of 96.75% by training the model using this method.

B. Confusion Matrix

A confusion matrix represents the performance of a classification model based on the ground truth values of the test dataset. The rows represent data instances of the predicted class and the columns represent data instances of the actual class. The figure 11: Confusion Matrix-1 represents the following:

0 : Downstairs
1 : Jogging
2 : Sitting
3 : Standing
4 : Upstairs
5 : Walking

Fig. 11.   Confusion Matrix - 1

IV. DISCUSSION

After computing the Confusion Matrix, we are able to draw various conclusions from the data in the dataset in relation to the model that we created. We have received 100% accuracy in the classes 'Jogging','Sitting' and 'Standing'. This shows that our model is able to classify co-ordinates relating to these classes effectively. However, our model is not without flaws. We can also deduce that our model is getting confused between co-ordinates relating to the classes 'Upstairs' and 'Downstairs'. If we observe Fig. 11, we can see that the True label 'Downstairs' denoted by Row 0 shows that only 64% of the test data has been correctly classified. The rest has

been classified as 'Upstairs'. This can be attributed to various reasons. The major reason being that our original dataset was highly unbalanced - the data was skewed towards the class 'Walking'. Thus in our endeavour to balance the dataset, we had to drop a lot of data amongst classes. We certainly would have had better results otherwise.

Fig. 12.   Confusion Matrix - 2

V. CONCLUSION

Human Activity Recognition has a vast array of applications including biometric signature, health and fitness monitoring, and eldercare. One of the main challenges in this field is the estimation of human poses. Traditionally, we have used a hand crafted model that requires specific initialisation and parameter estimation. But thanks to advances in Deep Learning, we are able to use Neural Networks for pose-tracking.

In conclusion, we have got a decent accuracy for this data. We have chosen StratifiedKFold for splitting the data because the class distribution in the dataset is preserved in the training and test splits, making it ideal for imbalanced datasets. In order to further increase the accuracy, we can play around with a few things. We can try training the model with more data or even try tuning the frame size and hop size. Furthermore, instead of using Deep Neural Networks for classification, we can use techniques such as Support Vector Machine (SVM), Random Forest (RF), and k-Nearest Neighbour.

REFERENCES

[1] Jeffrey W. Lockhart, Tony Pulickal, and Gary M. Weiss (2012). "Applications of Mobile Activity Recognition," Proceedings of the ACM UbiComp International Workshop on Situation, Activity, and Goal Awareness, Pittsburgh, PA.
[2] Gary M. Weiss and Jeffrey W. Lockhart (2012). "The Impact of Personalization on Smartphone-Based Activity Recognition," Proceedings of the AAAI-12 Workshop on Activity Context Representation: Techniques and Languages, Toronto, CA.
[3] Jennifer R. Kwapisz, Gary M. Weiss and Samuel A. Moore (2010). "Activity Recognition using Cell Phone Accelerometers," Proceedings of the Fourth International Workshop on Knowledge Discovery from Sensor Data (at KDD-10), Washington DC.

[4] Gupta, N., Gupta, S.K., Pathak, R.K. et al. Human activity recognition in artificial intelligence framework: a narrative review. Artif Intell Rev 55, 4755–4808 (2022).

[5] Jennifer R. Kwapisz, Gary M. Weiss and Samuel A. Moore (2010). Activity Recognition using Cell Phone Accelerometers, Proceedings of the Fourth International Workshop on Knowledge Discovery from Sensor Data (at KDD-10), Washington DC.

[6] Davide Anguita, Alessandro Ghio, Luca Oneto, Xavier Parra and Jorge L. Reyes-Ortiz. A Public Domain Dataset for Human Activity Recognition Using Smartphones. 21th European Symposium on Artificial Neural Networks, Computational Intelligence and Machine Learning, ESANN 2013.